

# 第二次实例分析

晏悦，任轶慧，徐逸斌，付娆

## 1. xv6中自旋锁 spinlock 的数据结构以及重要字段的作用

```
1 // Mutual exclusion lock.
2 struct spinlock {
3     uint locked;           // Is the lock held?
4
5     // For debugging:
6     char *name;            // Name of lock.
7     struct cpu *cpu;       // The cpu holding the lock.
8     uint pcs[10];          // The call stack (an array of program counters)
9                             // that locked the lock.
10 };
```

Locked:锁是否被持有?

1:被持有

0: 不被持有

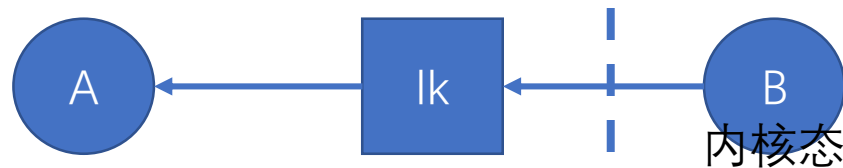
2. acquire 函数获取锁的时候, 为什么首先进行要执行 pushcli “关中断”操作?

```
20 // Acquire the lock.
21 // Loops (spins) until the lock is acquired.
22 // Holding a lock for a long time may cause
23 // other CPUs to waste time spinning to acquire it.
24 void
25 acquire(struct spinlock *lk)
26 {
27     pushcli(); // disable interrupts to avoid deadlock.
28     if(holding(lk))
29         panic("acquire");
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34
35     // Tell the C compiler and the processor to not move load
36     // past this point, to ensure that the critical section
37     // references happen after the lock is acquired.
38     __sync_synchronize();
39
40     // Record info about lock acquisition for debugging.
41     lk->cpu = mycpu();
42     getcallerpcs(&lk, lk->pcs);
43 }
```

理论上, 如果有中断程序可能获得锁lk, 则需要在acquire()开始处关中断。否则可能造成死锁。

实例:

- 应用程序A持有锁lk
- 应用程序A执行的过程中, 出现了中断
- 系统陷入内核, 执行中断处理程序B
- 中断处理程序B执行的过程中, 申请锁lk
- 但此时锁被应用程序A持有, 造成死锁



3. 内联汇编函数 `xchg` 返回 `lk->locked` 在 `xchg` 执行之前的值，考虑两个core上的分别有一个进程先后执行 `xchg`，返回值分别为0和1时，这两个进程分别发生了什么？

结合这个例子，说明 `xchg` 指令是原子指令的重要性

```
20 // Acquire the lock.
21 // Loops (spins) until the lock is acquired.
22 // Holding a lock for a long time may cause
23 // other CPUs to waste time spinning to acquire it.
24 void
25 acquire(struct spinlock *lk)
26 {
27     pushcli(); // disable interrupts to avoid deadlock.
28     if(holding(lk))
29         panic("acquire");
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34
35     // Tell the C compiler and the processor to not move loads
36     // past this point, to ensure that the critical section's
37     // references happen after the lock is acquired.
38     __sync_synchronize();
39
40     // Record info about lock acquisition for debugging.
41     lk->cpu = mycpu();
42     getcallerpcs(&lk, lk->pcs);
43 }
```

Xchg指令：交换arg1 和 arg2 中的内容，返回交换前arg1的值。

(1)锁被占用: `lk->locked == 1`, 返回值为1, 1写入`lk->locked`, 继续while循环内申请锁;  
(2)锁未被占用: `lk->locked == 0`, 返回值为0, 1写入`lk->locked`, 跳出while循环。

如果xchg为原子指令：

(1)A进程申请锁lk, `lk->locked==0`, 将`lk->locked`改写为1, 返回0;  
(2)B进程申请锁lk, `lk->locked==1`, 返回1;  
(3)进程A获得锁, 进程B继续申请。

3. 内联汇编函数 `xchg` 返回 `lk->locked` 在 `xchg` 执行之前的值，考虑两个core上的分别有一个进程先后执行 `xchg`，返回值分别为0和1时，这两个进程分别发生了什么？

结合这个例子，说明 `xchg` 指令是原子指令的重要性

```
20 // Acquire the lock.
21 // Loops (spins) until the lock is acquired.
22 // Holding a lock for a long time may cause
23 // other CPUs to waste time spinning to acquire it.
24 void
25 acquire(struct spinlock *lk)
26 {
27     pushcli(); // disable interrupts to avoid deadlock.
28     if(holding(lk))
29         panic("acquire");
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34
35     // Tell the C compiler and the processor to not move load
36     // past this point, to ensure that the critical section's
37     // references happen after the lock is acquired.
38     __sync_synchronize();
39
40     // Record info about lock acquisition for debugging.
41     lk->cpu = mycpu();
42     getcallerpcs(&lk, lk->pcs);
43 }
```

Xchg指令：交换arg1 和 arg2 中的内容，返回交换前arg1的值。

(1)锁被占用: `lk->locked == 1`, 返回值为1, 1写入`lk->locked`, 继续while循环内申请锁;  
(2)锁未被占用: `lk->locked == 0`, 返回值为0, 1写入`lk->locked`, 跳出while循环。

如果xchg不是原子指令：

Process A

Load &lk, %exa → 0  
Store &lk, %X ← 1

Process B

Load &lk, %exa → 0  
Store &lk, %X ← 1

两个进程都返回1，两个进程都获得锁。

4. 在编译后的xv6源码目录下执行 `objdump -d spinlock.o` , 观察 `acquire` 函数中, `xchg` 函数和它周围的while循环被编译成了什么汇编代码? 试将该汇编代码和C代码对应起来。

```
// Acquire the lock.
// Loops (spins) until the lock is acquired.
// Holding a lock for a long time may cause
// other CPUs to waste time spinning to acquire it.
void
acquire(struct spinlock *lk)
{
    pushcli(); // disable interrupts to avoid deadlock.
    if(holding(lk))
        panic("acquire");

    // The xchg is atomic.
    while(xchg(&lk->locked, 1) != 0)
        ;

    // Tell the C compiler and the processor to not move l
    // past this point, to ensure that the critical sectio
    // references happen after the lock is acquired.
    __sync_synchronize();

    // Record info about lock acquisition for debugging.
    lk->cpu = mycpu();
    getcallerpcs(&lk, lk->pcs);
}
```

```
138 00000170 <acquire>:
139 170: 55          push    %ebp
140 171: 89 e5       mov     %esp, %ebp
141 173: 53          push    %ebx
142 174: 83 ec 14    sub     $0x14, %esp
177: e8 fc ff ff ff call    178 <acquire+0x8>
17c: 8b 45 08     mov     0x8(%ebp), %eax
17f: 89 04 24     mov     %eax, (%esp)
182: e8 fc ff ff ff call    183 <acquire+0x13>
187: 85 c0       test    %eax, %eax
189: 75 3c       jne     1c7 <acquire+0x57>
18b: b9 01 00 00 00 mov     $0x1, %ecx
190: 8b 55 08     mov     0x8(%ebp), %edx
193: 89 c8       mov     %ecx, %eax
195: f0 87 02     lock xchg %eax, (%edx)
198: 85 c0       test    %eax, %eax
19a: 75 f4       jne     190 <acquire+0x20>
19c: f0 83 0c 24 00 lock orl $0x0, (%esp)
1a1: 8b 5d 08     mov     0x8(%ebp), %ebx
1a4: e8 fc ff ff ff call    1a5 <acquire+0x35>
1a9: 89 43 08     mov     %eax, 0x8(%ebx)
1ac: 8b 45 08     mov     0x8(%ebp), %eax
1af: 83 c0 0c     add     $0xc, %eax
1b2: 89 44 24 04 mov     %eax, 0x4(%esp)
1b6: 8d 45 08     lea     0x8(%ebp), %eax
1b9: 89 04 24     mov     %eax, (%esp)
1bc: e8 fc ff ff ff call    1bd <acquire+0x4d>
1c1: 83 c4 14     add     $0x14, %esp
1c4: 5b          pop     %ebx
1c5: 5d          pop     %ebp
1c6: c3          ret
```



## 5. 查阅内联汇编相关文档以及i386手册，结合objdump的输出，解释内联汇编函数 xchg (0568~0679)

### 内联汇编

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write op.
    asm volatile("lock; xchgl %0, %1" :
                  "+m" (*addr), "=a" (result) :
                  "1" (newval) :
                  "cc");
    return result;
}
```

防止GCC优化

```
138 00000170 <acquire>:
139 170: 55          push    %ebp
140 171: 89 e5       mov     %esp,%ebp
141 173: 53          push    %ebx
142 174: 83 ec 14    sub     $0x14,%esp
143 177: e8 fc ff ff call    178 <acquire+0x8>
144 17c: 8b 45 08    mov     0x8(%ebp),%eax
145 17f: 89 04 24    mov     %eax,(%esp)
146 182: e8 fc ff ff call    183 <acquire+0x13>
147 187: 85 c0       test    %eax,%eax
148 189: 75 3c       jne     1c7 <acquire+0x57>
149 18b: b9 01 00 00 mov     $0x1,%ecx
150 190: 8b 55 08    mov     0x8(%ebp),%edx
151 193: 89 c8       mov     %ecx,%eax
152 195: f0 87 02    lock xchgl %eax,(%edx)
153 198: 85 c0       test    %eax,%eax
154 19a: 75 f4       jne     190 <acquire+0x20>
155 19c: f0 83 0c 24 lock orl $0x0,(%esp)
156 1a1: 8b 5d 08    mov     0x8(%ebp),%ebx
157 1a4: e8 fc ff ff call    1a5 <acquire+0x35>
158 1a9: 89 43 08    mov     %eax,0x8(%ebx)
159 1ac: 8b 45 08    mov     0x8(%ebp),%eax
160 1af: 83 c0 0c    add     $0xc,%eax
161 1b2: 89 44 24 04 mov     %eax,0x4(%esp)
162 1b6: 8d 45 08    lea     0x8(%ebp),%eax
163 1b9: 89 04 24    mov     %eax,(%esp)
164 1bc: e8 fc ff ff call    1bd <acquire+0x4d>
165 1c1: 83 c4 14    add     $0x14,%esp
166 1c4: 5b         pop     %ebx
167 1c5: 5d         pop     %ebp
168 1c6: c3         ret
```

## 5. 查阅内联汇编相关文档以及i386手册，结合objdump的输出，解释内联汇编函数 xchg (0568~0679)

### 内联汇编

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write op.
    asm volatile("lock; xchgl %0, %1" :
                 "+m" (*addr), "=a" (result) :
                 "1" (newval) :
                 "cc");
    return result;
}
```

输入

```
138 00000170 <acquire>:
139 170: 55          push    %ebp
140 171: 89 e5      mov     %esp,%ebp
141 173: 53        push    %ebx
142 174: 83 ec 14   sub     $0x14,%esp
143 177: e8 fc ff ff ff call   178 <acquire+0x8>
144 17c: 8b 45 08   mov     0x8(%ebp),%eax
145 17f: 89 04 24   mov     %eax,(%esp)
146 182: e8 fc ff ff ff call   183 <acquire+0x13>
147 187: 85 c0      test    %eax,%eax
148 189: 75 3c     jne     1c7 <acquire+0x57>
149 18b: b9 01 00 00 00 mov     $0x1,%ecx
150 190: 8b 55 08   mov     0x8(%ebp),%edx
151 193: 89 c8      mov     %ecx,%eax
152 195: f0 87 02   lock xchgl %eax,(%edx)
153 198: 85 c0      test    %eax,%eax
154 19a: 75 f4     jne     190 <acquire+0x20>
155 19c: f0 83 0c 24 00 lock orl $0x0,(%esp)
156 1a1: 8b 5d 08   mov     0x8(%ebp),%ebx
157 1a4: e8 fc ff ff ff call   1a5 <acquire+0x35>
158 1a9: 89 43 08   mov     %eax,0x8(%ebx)
159 1ac: 8b 45 08   mov     0x8(%ebp),%eax
160 1af: 83 c0 0c   add     $0xc,%eax
161 1b2: 89 44 24 04 mov     %eax,0x4(%esp)
162 1b6: 8d 45 08   lea     0x8(%ebp),%eax
163 1b9: 89 04 24   mov     %eax,(%esp)
164 1bc: e8 fc ff ff ff call   1bd <acquire+0x4d>
165 1c1: 83 c4 14   add     $0x14,%esp
166 1c4: 5b        pop     %ebx
167 1c5: 5d        pop     %ebp
168 1c6: c3        ret
```



## 5. 查阅内联汇编相关文档以及i386手册，结合objdump的输出，解释内联汇编函数 xchg (0568~0679)

### 内联汇编

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write op.
    asm volatile("lock; xchgl %0, %1" :
                 "+m" (*addr), "=a" (result) :
                 "1" (newval) :
                 "cc");
    return result;
}
```

两个输出值  
'+'可读可输出， '='输出  
'm'内存， 'a'寄存器%eax

```
138 00000170 <acquire>:
139 170: 55          push    %ebp
140 171: 89 e5       mov     %esp, %ebp
141 173: 53          push    %ebx
142 174: 83 ec 14    sub     $0x14, %esp
143 177: e8 fc ff ff call    178 <acquire+0x8>
144 17c: 8b 45 08    mov     0x8(%ebp), %eax
145 17f: 89 04 24    mov     %eax, (%esp)
146 182: e8 fc ff ff call    183 <acquire+0x13>
147 187: 85 c0       test    %eax, %eax
148 189: 75 3c       jne     1c7 <acquire+0x57>
149 18b: b9 01 00 00 mov     $0x1, %ecx
150 190: 8b 55 08    mov     0x8(%ebp), %edx
151 193: 89 c8       mov     %ecx, %eax
152 195: f0 87 02    lock xchgl %eax, (%edx)
153 198: 85 c0       test    %eax, %eax
154 19a: 75 f4       jne     190 <acquire+0x20>
155 19c: f0 83 0c 24 lock orl $0x0, (%esp)
156 1a1: 8b 5d 08    mov     0x8(%ebp), %ebx
157 1a4: e8 fc ff ff call    1a5 <acquire+0x35>
158 1a9: 89 43 08    mov     %eax, 0x8(%ebx)
159 1ac: 8b 45 08    mov     0x8(%ebp), %eax
160 1af: 83 c0 0c    add     $0xc, %eax
161 1b2: 89 44 24 04 mov     %eax, 0x4(%esp)
162 1b6: 8d 45 08    lea     0x8(%ebp), %eax
163 1b9: 89 04 24    mov     %eax, (%esp)
164 1bc: e8 fc ff ff call    1bd <acquire+0x4d>
165 1c1: 83 c4 14    add     $0x14, %esp
166 1c4: 5b         pop     %ebx
167 1c5: 5d         pop     %ebp
168 1c6: c3         ret
```

## 5. 查阅内联汇编相关文档以及i386手册，结合objdump的输出，解释内联汇编函数 xchg (0568~0679)

### 内联汇编

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write op.
    asm volatile("lock; xchgl %0, %1" :
                 "+m" (*addr), "=a" (result) :
                 "l" (newval) :
                 "cc");
    return result;
}
```

“Lock”指令前缀

保证了这条指令对总线和缓存的独占权，  
保证了这条指令的执行过程中不会有其他CPU  
或同CPU内的指令访问缓存和内存

```
138 00000170 <acquire>:
139 170: 55          push    %ebp
140 171: 89 e5       mov     %esp, %ebp
141 173: 53          push    %ebx
142 174: 83 ec 14    sub     $0x14, %esp
143 177: e8 fc ff ff call    178 <acquire+0x8>
144 17c: 8b 45 08     mov     0x8(%ebp), %eax
145 17f: 89 04 24     mov     %eax, (%esp)
146 182: e8 fc ff ff call    183 <acquire+0x13>
147 187: 85 c0       test    %eax, %eax
148 189: 75 3c       jne     1c7 <acquire+0x57>
149 18b: b9 01 00 00 mov     $0x1, %ecx
150 190: 8b 55 08     mov     0x8(%ebp), %edx
151 193: 89 c8       mov     %ecx, %eax
152 195: f0 87 02    lock xchgl %eax, (%edx)
153 198: 85 c0       test    %eax, %eax
154 19a: 75 f4       jne     190 <acquire+0x20>
155 19c: f0 83 0c 24 lock orl $0x0, (%esp)
156 1a1: 8b 5d 08     mov     0x8(%ebp), %ebx
157 1a4: e8 fc ff ff call    1a5 <acquire+0x35>
158 1a9: 89 43 08     mov     %eax, 0x8(%ebx)
159 1ac: 8b 45 08     mov     0x8(%ebp), %eax
160 1af: 83 c0 0c     add     $0xc, %eax
161 1b2: 89 44 24 04 mov     %eax, 0x4(%esp)
162 1b6: 8d 45 08     lea     0x8(%ebp), %eax
163 1b9: 89 04 24     mov     %eax, (%esp)
164 1bc: e8 fc ff ff call    1bd <acquire+0x4d>
165 1c1: 83 c4 14     add     $0x14, %esp
166 1c4: 5b          pop     %ebx
167 1c5: 5d          pop     %ebp
168 1c6: c3          ret
```

## 6. acquire 和 release 函数中的 \_\_sync\_synchronize() 语句的作用是什么？在汇编程序中是如何体现的？

- 设置一个内存屏障（memory barrier），防止编译器reorder的时候把临界区的访存指令移到锁操作之前
- 该语句被编译为lock orl \$0x0, (%esp)指令，lock指令前缀...

```
45 // Release the lock.
46 void
47 release(struct spinlock *lk)
48 {
49     if(!holding(lk))
50         panic("release");
51
52     lk->pcs[0] = 0;
53     lk->cpu = 0;
54
55     // Tell the C compiler and the processor to not move loads or stores
56     // past this point, to ensure that all the stores in the critical
57     // section are visible to other cores before the lock is released.
58     // Both the C compiler and the hardware may re-order loads and
59     // stores; __sync_synchronize() tells them both not to.
60     __sync_synchronize();
61
62     // Release the lock, equivalent to lk->locked = 0.
63     // This code can't use a C assignment, since it might
64     // not be atomic. A real OS would use C atomics here.
65     asm volatile("movl $0, %0" : "+m" (lk->locked) : );
66
67     popcli();
68 }
```

```
195 00000210 <release>:
196 210: 55                push    %ebp
197 211: 89 e5            mov     %esp,%ebp
198 213: 53              push    %ebx
199 214: 83 ec 10        sub     $0x10,%esp
200 217: 8b 5d 08        mov     0x8(%ebp),%ebx
201 21a: 53              push    %ebx
202 21b: e8 fc ff ff ff  call    21c <release+0xc>
203 220: 83 c4 10        add     $0x10,%esp
204 223: 85 c0            test    %eax,%eax
205 225: 74 22            je      249 <release+0x39>
206 227: c7 43 0c 00 00 00 00  movl    $0x0,0xc(%ebx)
207 22e: c7 43 08 00 00 00 00  movl    $0x0,0x8(%ebx)
208 235: f0 83 0c 24 00  lock orl $0x0, (%esp)
209 23a: c7 03 00 00 00 00  movl    $0x0, (%ebx)
210 240: 8b 5d fc        mov     -0x4(%ebp),%ebx
211 243: c9              leave
212 244: e9 77 fe ff ff  jmp     c0 <popcli>
213 249: 83 ec 0c        sub     $0xc,%esp
214 24c: 68 26 00 00 00  push    $0x26
215 251: e8 fc ff ff ff  call    252 <release+0x42>
216
```

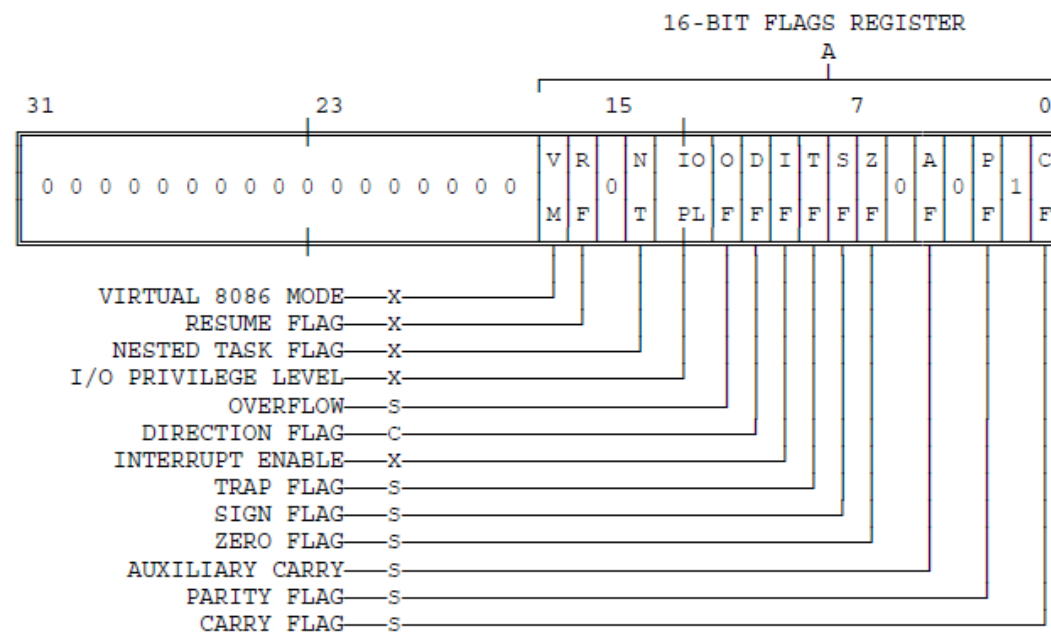
## 7. 查阅i386手册，了解pushcli函数中的EFLAGS和FL\_IF分别是什么？ sti和cli内联汇编函数是如何影响前者的？

- EFLAGS: 32位的标志位寄存器
- FL\_IF: xv6自定义的宏，与EFLAGS相与可以得到中断位IF

```
1 // This file contains definitions for the
2 // x86 memory management unit (MMU).
3
4 // Eflags register
5 #define FL_IF      0x00000200    // Interrupt Enable
6
```

```
104 void
105 pushcli(void)
106 {
107     int eflags;
108
109     eflags = readeflags();
110     cli();
111     if(mycpu()->ncli == 0)
112         mycpu()->intena = eflags & FL_IF;
113     mycpu()->ncli += 1;
114 }
```

Figure 2-8. EFLAGS Register



S = STATUS FLAG, C = CONTROL FLAG, X = SYSTEM FLAG

NOTE: 0 OR 1 INDICATES INTEL RESERVED. DO NOT DEFINE

7. 查阅i386手册，了解pushcli函数中的EFLAGS和FL\_IF分别是什么？ sti和cli内联汇编函数是如何影响前者的？

- STI：设置 EFLAGS 寄存器中的中断标志 (IF)。设置 IF 标志之后，处理器在执行下一个指令之后，开始响应可屏蔽的外部中断。（例如STI+RET）
- CLI：如果CPL小于等于IOPL，清除EFLAGS寄存器中的IF标志，否则不清除，EFLAGS的其他标志位不会受到影响。

## 8. mycpu()->intena变量的作用是什么?

- 表示在调用pushcli函数之前是否允许中断
- pushcli和popcli的调用过程...

```
void
pushcli(void)
{
    int eflags;

    eflags = readeflags();
    cli();
    if(mycpu()->ncli == 0)
        mycpu()->intena = eflags & FL_IF;
    mycpu()->ncli += 1;
}
```

```
// Per-CPU state
struct cpu {
    uchar apicid;                // Local APIC ID
    struct context *scheduler;   // swtch() here to enter scheduler
    struct taskstate ts;         // Used by x86 to find stack for interrupt
    struct segdesc gdt[NSEGS];   // x86 global descriptor table
    volatile uint started;       // Has the CPU started?
    int ncli;                    // Depth of pushcli nesting.
    int intena;                  // Were interrupts enabled before pushcli?
    struct proc *proc;           // The process running on this cpu or null
};
```

```
void
popcli(void)
{
    if(readeflags() & FL_IF)
        panic("popcli - interruptible");
    if(--mycpu()->ncli < 0)
        panic("popcli");
    if(mycpu()->ncli == 0 && mycpu()->intena)
        sti();
}
```



## Linux 锁相关代码分析(x86版本)

- 1. x86实现原子指令
- 方法: x86提供了附加的lock前缀, 使带**lock前缀**的读修改写指令能原子性执行。
- 实现原理: 带lock前缀的指令在操作时会**锁住总线**, 使自身的执行即使在多处理器间也是原子性执行的。
- API使用: 原子性操作是线程间同步的基础, linux专门定义了一种只进行原子操作的类型atomic\_t, 并提供相关的原子读写调用API。

# atomic\_t数据结构与API

- typedef struct {
- volatile int counter;
- } atomic\_t;
- 原子类型其实是int类型，只是禁止寄存器对其暂存
- static inline int atomic\_read(const atomic\_t \*v)
- {
- return v->counter;
- }
- static inline void atomic\_set(atomic\_t \*v, int i)
- {
- v->counter = i;
- }
- 单独读或者写, 在x86下是原子性的

# 原子加, 原子减

- static inline void atomic\_add(int i, atomic\_t \*v)
- {
- asm volatile(LOCK\_PREFIX "addl %1,%0"
- : "+m" (v->counter)
- : "ir" (i));
- }
- 
- static inline void atomic\_sub(int i, atomic\_t \*v)
- {
- asm volatile(LOCK\_PREFIX "subl %1,%0"
- : "+m" (v->counter)
- : "ir" (i));
- }

**LOCK\_PREFIX**  
原子指令前缀

counter值加上i 再存回  
counter

# linux FIFO ticket-based spinlock

- 1. 普通自旋锁的缺陷: 由于无序竞争的本质特点, 内核执行线程无法保证何时可以取到锁, 某些执行线程可能需要等待很长时间, 导致“不公平”问题的产生。
- 2. 改进方法: 先申请锁的排一个号码, 当锁释放时候, 然后按照号码来调度号码小的任务(类似去银行办理服务, 假设只有一个窗口, 每个人有排队号码)
- 3. 数据结构: next 与 owner 域, 均初始化为0
- next: 等待服务的任务数目,
- owner: 已经服务完成的任务数目

next: 等待服务的任务数目,  
owner: 已经服务完成的任务数目

```
acquire{  
    tmp = fetch_and_add(next)  
    while( tmp != owner )  
        ;  
}
```

next++,  
返回原来值  
新来了一个服务

当前完成任务不等于等待服  
务任务数, 忙等待(直到排队  
号到了)

```
release{  
    atomic_add(owner);  
}
```

服务完成任务数++