

# XV6 条件变量

报告人：高亦钊

小组成员：高亦钊，索玉玺，王子铭，范楷元

# Background

考虑一个最简单的例子，一对发送者和接受者（消费者和生产者）

Send 进程将一个数据的指针发送给Recv

```
100 struct q {  
101     void *ptr;  
102 };  
103  
104 void*  
105 send(struct q *q, void *p)  
106 {  
107     while(q->ptr != 0)  
108         ;  
109     q->ptr = p;  
110 }  
111
```

```
112 void*  
113 recv(struct q *q)  
114 {  
115     void *p;  
116  
117     while((p = q->ptr) == 0)  
118         ;  
119     q->ptr = 0;  
120     return p;  
121 }
```

# Ver 1.0

优点: 没有数据丢失

缺点: 如果很久发一次数据, recv一直在判断是否有新数据, 浪费CPU资源

```
100 struct q {  
101     void *ptr;  
102 };  
103  
104 void*  
105 send(struct q *q, void *p)  
106 {  
107     while(q->ptr != 0)  
108         ;  
109     q->ptr = p;  
110 }  
111
```

```
112 void*  
113 recv(struct q *q)  
114 {  
115     void *p;  
116  
117     while((p = q->ptr) == 0)  
118         ;  
119     q->ptr = 0;  
120     return p;  
121 }
```

## Ver 2.0

假象我们有一组函数wakeup和sleep,让recv函数在没有新数据的时候放弃CPU资源, 当有新数据的时候send将recv唤醒

新的问题: lost wake-up

```
201 void*
202 send(struct q *q, void *p)
203 {
204     while(q->ptr != 0)
205         ;
206     q->ptr = p;
207     wakeup(q); /* wake recv */
208 }
209
```

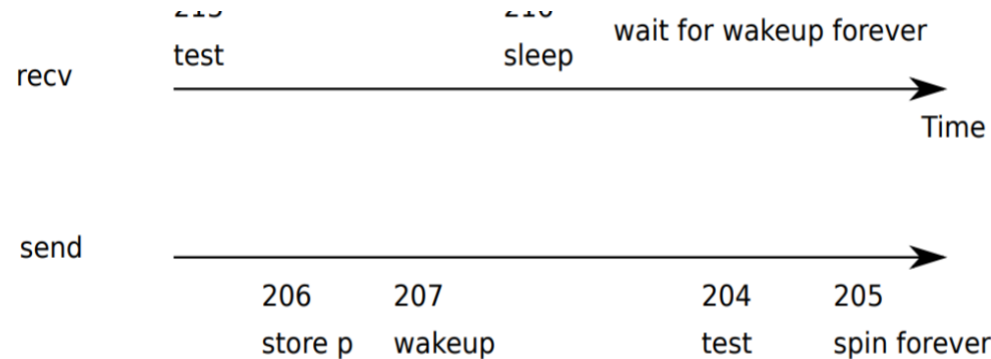
```
210 void*
211 recv(struct q *q)
212 {
213     void *p;
214
215     while((p = q->ptr) == 0)
216         sleep(q);
217     q->ptr = 0;
218     return p;
219 }
```

# Lost wake-up

如果recv函数在结束判断决定执行sleep,但是sleep函数还没有彻底执行完时, send将新数据写入并且调用wakeup,此时发现并没有进程在sleep

```
201 void*
202 send(struct q *q, void *p)
203 {
204     while(q->ptr != 0)
205         ;
206     q->ptr = p;
207     wakeup(q); /* wake recv */
208 }
209
```

```
210 void*
211 recv(struct q *q)
212 {
213     void *p;
214
215     while((p = q->ptr) == 0)
216         sleep(q);
217     q->ptr = 0;
218     return p;
219 }
```



原因总结: 没有保证sleep成功前没有其他进程来发送新的数据

# Ver 3.0: 加锁（在进入临界区的时候加锁）

## 新的问题：产生死锁（sleep结束后锁没有释放）

```
300 struct q {
301     struct spinlock lock;
302     void *ptr;
303 };
304
305 void*
306 send(struct q *q, void *p)
307 {
308     acquire(&q->lock);
309     while(q->ptr != 0)
310         ;
311     q->ptr = p;
312     wakeup(q);
313     release(&q->lock);
314 }
315
```

```
316 void*
317 recv(struct q *q)
318 {
319     void *p;
320
321     acquire(&q->lock);
322     while((p = q->ptr) == 0)
323         sleep(q);
324     q->ptr = 0;
325     release(&q->lock);
326     return p;
327 }
```

Ver 4.0:我们需要让sleep函数在结束的时候将锁释放，因此需要将锁作为参数，让sleep成功之后释放锁

```
400 struct q {
401     struct spinlock lock;
402     void *ptr;
403 };
404
405 void*
406 send(struct q *q, void *p)
407 {
408     acquire(&q->lock);
409     while(q->ptr != 0)
410         ;
411     q->ptr = p;
412     wakeup(q);
413     release(&q->lock);
414 }
```

```
415
416 void*
417 recv(struct q *q)
418 {
419     void *p;
420
421     acquire(&q->lock);
422     while((p = q->ptr) == 0)
423         sleep(q, &q->lock);
424     q->ptr = 0;
425     release(&q->lock);
426     return p;
427 }
```



## Sleep

参数: \*chan

(Callers of sleep and wakeup can use any mutually convenient number as the channel. Xv6 often uses the address of a kernel data structure involved in the waiting.)

参数: \*lk (问题2)

为什么需要锁前面已经提及:

1.锁用来保证lost wake-up, 让sleep能安全的进入睡眠

2.Sleep函数需要将锁释放, 否则这把锁不会被释放, wakeup函数无法执行

```
415 // Atomically release lock and sleep on chan.
416 // Reacquires lock when awakened.
417 void
418 sleep(void *chan, struct spinlock *lk)
419 {
420     struct proc *p = myproc();
421
422     if(p == 0)
423         panic("sleep");
424
425     if(lk == 0)
426         panic("sleep without lk");
427
428     // Must acquire ptable.lock in order to
429     // change p->state and then call sched.
430     // Once we hold ptable.lock, we can be
431     // guaranteed that we won't miss any wakeup
432     // (wakeup runs with ptable.lock locked),
433     // so it's okay to release lk.
434     if(lk != &ptable.lock){ //DOC: sleeplock0
435         acquire(&ptable.lock); //DOC: sleeplock1
436         release(lk);
437     }
438     // Go to sleep.
439     p->chan = chan;
440     p->state = SLEEPING;
441
442     sched();
```

```
...
444 // Tidy up.
445 p->chan = 0;
446
447 // Reacquire original lock.
448 if(lk != &ptable.lock){ //DOC: sleeplock2
449     release(&ptable.lock);
450     acquire(lk);
451 }
452 }
453
454 //PAGEBREAK!
455 // Wake up all processes sleeping on chan.
456 // The ptable lock must be held.
457 static void
458 wakeup1(void *chan)
459 {
460     struct proc *p;
461
462     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
463         if(p->state == SLEEPING && p->chan == chan)
464             p->state = RUNNABLE;
465 }
```



问题3: 为什么 sleep (以及 wakeup ) 要使用 ptable.lock ? xv6如何通过锁的使用来解决lost wake-up问题?

为什么要使用ptable.lock:  
因为涉及到线程状态的变化

```
415 // Atomically release lock and sleep on chan.
416 // Reacquires lock when awakened.
417 void
418 sleep(void *chan, struct spinlock *lk)
419 {
420     struct proc *p = myproc();
421
422     if(p == 0)
423         panic("sleep");
424
425     if(lk == 0)
426         panic("sleep without lk");
427
428     // Must acquire ptable.lock in order to
429     // change p->state and then call sched.
430     // Once we hold ptable.lock, we can be
431     // guaranteed that we won't miss any wakeup
432     // (wakeup runs with ptable.lock locked),
433     // so it's okay to release lk.
434     if(lk != &ptable.lock){ //DOC: sleeplock0
435         acquire(&ptable.lock); //DOC: sleeplock1
436         release(lk);
437     }
438     // Go to sleep.
439     p->chan = chan;
440     p->state = SLEEPING;
441
442     sched();
```

```
444 // Tidy up.
445 p->chan = 0;
446
447 // Reacquire original lock.
448 if(lk != &ptable.lock){ //DOC: sleeplock2
449     release(&ptable.lock);
450     acquire(lk);
451 }
452 }
453
454 // Wake up all processes sleeping on chan.
455 void
456 wakeup(void *chan)
457 {
458     acquire(&ptable.lock);
459     wakeup1(chan);
460     release(&ptable.lock);
461 }
462
463 //PAGEBREAK!
464 // Wake up all processes sleeping on chan.
465 // The ptable lock must be held.
466 static void
467 wakeup1(void *chan)
468 {
469     struct proc *p;
470
471     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
472         if(p->state == SLEEPING && p->chan == chan)
473             p->state = RUNNABLE;
474 }
```

问题4: sleep 函数中, 2890行if语句的作用是什么?  
如果lk是ptable.lock, 防止再申请一次, 并且保证sleep不会当成正常的一个lock并释放掉

```
415 // Atomically release lock and sleep on chan.
416 // Reacquires lock when awakened.
417 void
418 sleep(void *chan, struct spinlock *lk)
419 {
420     struct proc *p = myproc();
421
422     if(p == 0)
423         panic("sleep");
424
425     if(lk == 0)
426         panic("sleep without lk");
427
428     // Must acquire ptable.lock in order to
429     // change p->state and then call sched.
430     // Once we hold ptable.lock, we can be
431     // guaranteed that we won't miss any wakeup
432     // (wakeup runs with ptable.lock locked),
433     // so it's okay to release lk.
434     if(lk != &ptable.lock){ //DOC: sleeplock0
435         acquire(&ptable.lock); //DOC: sleeplock1
436         release(lk);
437     }
438     // Go to sleep.
439     p->chan = chan;
440     p->state = SLEEPING;
441
442     sched();
```

```
...
444 // Tidy up.
445 p->chan = 0;
446
447 // Reacquire original lock.
448 if(lk != &ptable.lock){ //DOC: sleeplock2
449     release(&ptable.lock);
450     acquire(lk);
451 }
452 }
453
454 // Wake up all processes sleeping on chan.
455 void
456 wakeup(void *chan)
457 {
458     acquire(&ptable.lock);
459     wakeup1(chan);
460     release(&ptable.lock);
461 }
462
463 //PAGEBREAK!
464 // Wake up all processes sleeping on chan.
465 // The ptable lock must be held.
466 static void
467 wakeup1(void *chan)
468 {
469     struct proc *p;
470
471     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
472         if(p->state == SLEEPING && p->chan == chan)
473             p->state = RUNNABLE;
474 }
```

问题5: sleep 函数中, 2891行和2892行能不能交换顺序? 为什么?

不能! 在现在这种情况下sleep在释放lk前会获得ptable.lock, 同时拥有两把锁, 保证没有进程能执行wakeup

如果被调换位置, 先释放lk,再申请ptable.lock, 就会存在一个时刻sleep还没拿到ptable.lock, 但是释放了lk, 这时如果lk和ptable.lock都被一个进程得到, 那么就可以执行wakeup,造成lost wakeup

```
415 // Atomically release lock and sleep on chan.
416 // Reacquires lock when awakened.
417 void
418 sleep(void *chan, struct spinlock *lk)
419 {
420     struct proc *p = myproc();
421
422     if(p == 0)
423         panic("sleep");
424
425     if(lk == 0)
426         panic("sleep without lk");
427
428     // Must acquire ptable.lock in order to
429     // change p->state and then call sched.
430     // Once we hold ptable.lock, we can be
431     // guaranteed that we won't miss any wakeup
432     // (wakeup runs with ptable.lock locked),
433     // so it's okay to release lk.
434     if(lk != &ptable.lock){ //DOC: sleeplock0
435         acquire(&ptable.lock); //DOC: sleeplock1
436         release(lk);
437     }
438     // Go to sleep.
439     p->chan = chan;
440     p->state = SLEEPING;
441
442     sched();
```

```
444 // Tidy up.
445 p->chan = 0;
446
447 // Reacquire original lock.
448 if(lk != &ptable.lock){ //DOC: sleeplock2
449     release(&ptable.lock);
450     acquire(lk);
451 }
452 }
453
454 // Wake up all processes sleeping on chan.
455 void
456 wakeup(void *chan)
457 {
458     acquire(&ptable.lock);
459     wakeup1(chan);
460     release(&ptable.lock);
461 }
462
463 //PAGEBREAK!
464 // Wake up all processes sleeping on chan.
465 // The ptable lock must be held.
466 static void
467 wakeup1(void *chan)
468 {
469     struct proc *p;
470
471     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
472         if(p->state == SLEEPING && p->chan == chan)
473             p->state = RUNNABLE;
474 }
```

问题6: 阐述 sleep 函数执行时, 进程是如何转入睡眠态, 又转入就绪态和运行态, 并继续执行sleep 的?

进入就绪态:

1.将p->state 设为sleeping ;

p->chan 设为chan标记睡再哪个通道上

2.调用sche(), switch context, 进入调度程序, 交出CPU资源

```
415 // Atomically release lock and sleep on chan.
416 // Reacquires lock when awakened.
417 void
418 sleep(void *chan, struct spinlock *lk)
419 {
420     struct proc *p = myproc();
421
422     if(p == 0)
423         panic("sleep");
424
425     if(lk == 0)
426         panic("sleep without lk");
427
428     // Must acquire ptable.lock in order to
429     // change p->state and then call sched.
430     // Once we hold ptable.lock, we can be
431     // guaranteed that we won't miss any wakeup
432     // (wakeup runs with ptable.lock locked),
433     // so it's okay to release lk.
434     if(lk != &ptable.lock){ //DOC: sleeplock0
435         acquire(&ptable.lock); //DOC: sleeplock1
436         release(lk);
437     }
438     // Go to sleep.
439     p->chan = chan;
440     p->state = SLEEPING;
441
442     sched();
```

```
444 // Tidy up.
445 p->chan = 0;
446
447 // Reacquire original lock.
448 if(lk != &ptable.lock){ //DOC: sleeplock2
449     release(&ptable.lock);
450     acquire(lk);
451 }
452 }
453
454 // Wake up all processes sleeping on chan.
455 void
456 wakeup(void *chan)
457 {
458     acquire(&ptable.lock);
459     wakeup1(chan);
460     release(&ptable.lock);
461 }
462
463 //PAGEBREAK!
464 // Wake up all processes sleeping on chan.
465 // The ptable lock must be held.
466 static void
467 wakeup1(void *chan)
468 {
469     struct proc *p;
470
471     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
472         if(p->state == SLEEPING && p->chan == chan)
473             p->state = RUNNABLE;
474 }
```

问题6: 阐述 sleep 函数执行时, 进程是如何转入睡眠态, 又转入就绪态和运行态, 并继续执行sleep 的?

如何进入就绪态和运行态:

- 1.wakeup1函数将睡眠上所有的进程都设为runnable
- 2.等CPU接下来进行调度的时候就能将某个进程唤醒, 变为running, 并回到sleep调用sched时的context

Context switch from 380 to 346 and from 346 to 380

```
365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()->ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags() & FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()->intena;
380     swtch(&p->context, mycpu()->scheduler)
381     mycpu()->intena = intena;
382 }
```

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchuvm(p);
344             p->state = RUNNING;
345
346             swtch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354
355     }
356 }
```

问题7: xv6的 wakeup 操作, 为什么要拆分成 wakeup 和 wakeup1 两个函数, 请举例说明。

如果调用wakeup的进程已经拥有的ptable.lock, 重新申请的时候就会出问题。

例子exit函数, 已经拥有了ptable.lock, 这么用时因为之后还要更改其他process的状态, 省的再acquire一次

```
467 // Wake up all processes sleeping on chan.
468 void
469 wakeup(void *chan)
470 {
471     acquire(&ptable.lock);
472     wakeup1(chan);
473     release(&ptable.lock);
474 }
475
454 //PAGEBREAK!
455 // Wake up all processes sleeping on chan.
456 // The ptable lock must be held.
457 static void
458 wakeup1(void *chan)
459 {
460     struct proc *p;
461
462     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
463         if(p->state == SLEEPING && p->chan == chan)
464             p->state = RUNNABLE;
465 }
172 // Exit the current process. Does not return.
173 // An exited process remains in the zombie state
174 // until its parent calls wait() to find out it exited.
```

```
175 void
176 exit(void)
177 {
178     struct proc *p;
179     int fd;
180
181     if(proc == initproc){
182         panic("init exiting");
183     }
184
185     // Close all open files.
186     for(fd = 0; fd < NOFILE; fd++){
187         if(proc->ofile[fd]){
188             fileclose(proc->ofile[fd]);
189             proc->ofile[fd] = 0;
190         }
191     }
192
193     iput(proc->cwd);
194     proc->cwd = 0;
195
196     acquire(&ptable.lock);
197
198     // Parent might be sleeping in wait().
199     wakeup1(proc->parent);
200
201     // Pass abandoned children to init.
202     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
203         if(p->parent == proc){
204             p->parent = initproc;
205             if(p->state == ZOMBIE)
206                 wakeup1(initproc);
207         }
208     }
209 }
```



问题8: 假设 wakeup 操作唤醒了多个等待相同 channel 的进程, 此时这多个进程会如何执行? xv6的 wakeup 是否符合Mesa semantics?

1.会全部唤醒

2.如果数据已经被消费了, 被唤醒的程序就会发现没有数据可以处理, 由于scheduler进行上下文切换回到的是sleep调用sched的位置, 当sleep退出后, 会再次检测是否满足条件, 如果不满足进入睡眠, 因此保证了不发生错误。

Spurious wakeup

符合Mesa semantics(条件不为真, 重新判断一次)

```
while((p = q->ptr) == 0)
    sleep(q, &q->lock);
```

```
365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()->ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags() & FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()->intena;
380     swtch(&p->context, mycpu()->scheduler);
381     mycpu()->intena = intena;
382 }
---
```

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchvm(p);
344             p->state = RUNNING;
345
346             swtch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354
355     }
356 }
```



问题9: wakeup 时如果没有 sleeping 的进程, wakeup 会阻塞吗?

不会

```
467 // Wake up all processes sleeping on chan.
468 void
469 wakeup(void *chan)
470 {
471     acquire(&ptable.lock);
472     wakeup1(chan);
473     release(&ptable.lock);
474 }
475
454 //PAGEBREAK!
455 // Wake up all processes sleeping on chan.
456 // The ptable lock must be held.
457 static void
458 wakeup1(void *chan)
459 {
460     struct proc *p;
461
462     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
463         if(p->state == SLEEPING && p->chan == chan)
464             p->state = RUNNABLE;
465 }
```

问题1: 请查阅资料, 给出 pthread 提供的条件变量操作, 并与xv6提供的 sleep & wakeup 操作比较。

1. 初始化条件变量 pthread\_cond\_init

```
int pthread_cond_init(pthread_cond_t *cv, const
pthread_condattr_t *cattr);
```

2. 阻塞 pthread\_cond\_wait

```
int pthread_cond_wait(pthread_cond_t *cv,
pthread_mutex_t *mutex);
```

3. 解除条件变量上的阻塞 pthread\_cond\_signal

```
int pthread_cond_signal(pthread_cond_t *cv);
```

4. 定时阻塞 pthread\_cond\_timewait

```
int pthread_cond_timedwait(pthread_cond_t *cv,
pthread_mutex_t *mp, const struct timespec * abstime);
```

5. 释放阻塞的所有进程 pthread\_cond\_broadcast

```
int pthread_cond_broadcast(pthread_cond_t *cv);
```

6. 释放条件变量 pthread\_cond\_destroy

```
int pthread_cond_destroy(pthread_cond_t *cv);
```

对比:

1.pthread条件变量有自己的数据结构, xv6是自行定义的, 只需要满足有锁。

2.pthread 条件变量分为signal和broadcast,但是XV6只有“broadcast”. 使用signal的时候只唤醒一个进程, 哪一个又调度方法决定

3. Pthread还支持定时阻塞和释放条件变量

```
typedef struct {
    struct pthread_queue queue;
    int flags;
    int waiters;
    pthread_mutex_t *mutex;
} pthread_cond_t;
```

## 1. Condition Variable Data Structure

```
typedef struct {  
    struct pthread_queue queue;  
    int flags;  
    int waiters;  
    pthread_mutex_t *mutex;  
} pthread_cond_t;
```

## 2. Initialization

PTHREAD\_COND\_INITIALIZER

```
int pthread_cond_init(pthread_cond_t *cv, const pthread_condattr_t *cattr);
```

## 3. Wait

```
int pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex);
```

## 4. Signal

```
int pthread_cond_signal(pthread_cond_t *cv);
```

```
int pthread_cond_timedwait(pthread_cond_t *cv, pthread_mutex_t *mp, const struct timespec * abstime);
```

```
int pthread_cond_broadcast(pthread_cond_t *cv);
```

```
int pthread_cond_destroy(pthread_cond_t *cv);
```

# Pthread提供的条件变量操作

1. 初始化条件变量 pthread\_cond\_init

```
int pthread_cond_init(pthread_cond_t *cv, const pthread_condattr_t *cattr);
```

2. 阻塞 pthread\_cond\_wait

```
int pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex);
```

3. 解除条件变量上的阻塞 pthread\_cond\_signal

```
int pthread_cond_signal(pthread_cond_t *cv);
```

4. 定时阻塞 pthread\_cond\_timewait

```
int pthread_cond_timedwait(pthread_cond_t *cv, pthread_mutex_t *mp, const struct timespec * abstime);
```

5. 释放阻塞的所有进程 pthread\_cond\_broadcast

```
int pthread_cond_broadcast(pthread_cond_t *cv);
```

6. 释放条件变量 pthread\_cond\_destroy

```
int pthread_cond_destroy(pthread_cond_t *cv);
```

# Linux 信号量

## 问题1: Linux中信号量的数据结构及其对应的PV操作

### 1、Linux中信号量的数据结构:

结构体名:

struct semaphore

自旋锁: raw\_spinlock\_t lock;

信号量长度: unsigned int count;

等待进程的链表: struct list\_head wait\_list;

### 2、对应的PV操作:

#### (1) P操作

上锁: raw\_spin\_lock\_irqsave(&sem->lock, flags);

信号量大于0, 获得信号量: if (likely(sem->count > 0))

sem->count--;

信号量小于等于0, 休眠: else

\_\_down(sem);

释放锁: raw\_spin\_unlock\_irqrestore(&sem->lock, flags);

```
struct semaphore {
    raw_spinlock_t    lock;
    unsigned int      count;
    struct list_head   wait_list;
};

void down(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        __down(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}

EXPORT_SYMBOL(down);
```

## 问题1: Linux中信号量的数据结构及其对应的PV操作

- 2、对应的PV操作
- (2) V操作
- 上锁: `raw_spin_lock_irqsave(&sem->lock, flags);`
- 等待队列为空, 信号量加1: `if (likely(list_empty(&sem->wait_list)))`
- `sem->count++;`
- 等待队列非空, 唤醒: `else`
- `__up(sem);`
- 释放锁: `raw_spin_unlock_irqrestore(&sem->lock, flags);`

```
void up(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(list_empty(&sem->wait_list)))
        sem->count++;
    else
        __up(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}
EXPORT_SYMBOL(up);
```



## 问题2：说明Linux扩展的各类down操作的用途

- 1、down：在此函数中首先进行信号量资源数的查看，如果信号量数据(count)不为0，则将其减1，并返回，调用成功；否则调用\_\_down进行等待，调用者进行睡眠；该函数的调用不允许中断。
- 2、down\_interruptible：该函数功能和down类似，不同之处为，down不会被信号（signal）打断，但down\_interruptible能被信号打断，因此该函数有返回值来区分是正常返回还是被信号中断，如果返回0，表示获得信号量正常返回，如果被信号打断，返回-EINTR。
- 3、down\_killable：down\_killable与down\_interruptible相近，最终传入的\_\_down\_common的实参有所不同（TASK\_KILLABLE和TASK\_INTERRUPTIBLE），睡眠的进程可以因为受到致命信号而被唤醒，中断获取信号量的操作。

```
void down(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        __down(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}

EXPORT_SYMBOL(down);
```

```
int down_interruptible(struct semaphore *sem)
{
    unsigned long flags;
    int result = 0;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        result = __down_interruptible(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);

    return result;
}

EXPORT_SYMBOL(down_interruptible);
```

```
int down_killable(struct semaphore *sem)
{
    unsigned long flags;
    int result = 0;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        result = __down_killable(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);

    return result;
}
```

## 问题2：说明Linux扩展的各类down操作的用途

- 4、down\_trylock：试图获取资源，若无法获得则直接返回1而不睡眠。返回0则表示获取到了资源
- 5、down\_timeout：可以自定义超时时间，也就是如果在超时间内不能得到资源，调用者会因为超时而自行唤醒。其实现过程如下，请注意超时参数的传入。其中TASK\_UNINTERRUPTIBLE，如down一样，不可被致死信号和信号打断。

```
int down_trylock(struct semaphore *sem)
{
    unsigned long flags;
    int count;

    raw_spin_lock_irqsave(&sem->lock, flags);
    count = sem->count - 1;
    if (likely(count >= 0))
        sem->count = count;
    raw_spin_unlock_irqrestore(&sem->lock, flags);

    return (count < 0);
}
EXPORT_SYMBOL(down_trylock);
```

```
int down_timeout(struct semaphore *sem, long timeout)
{
    unsigned long flags;
    int result = 0;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        result = __down_timeout(sem, timeout);
    raw_spin_unlock_irqrestore(&sem->lock, flags);

    return result;
}
EXPORT_SYMBOL(down_timeout);
```

### 问题三：简要分析down和up操作的实现

- **down操作的实现**

上锁：raw\_spin\_lock\_irqsave(&sem->lock, flags);

信号量大于0，获得信号量：if (likely(sem->count > 0))  
sem->count--;

信号量小于等于0，休眠：else  
\_\_down(sem);

释放锁：raw\_spin\_unlock\_irqrestore(&sem->lock, flags)

```
void down(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        __down(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}
EXPORT_SYMBOL(down);
```

### 问题三：简要分析down和up操作的实现

- **down操作的实现**
- 2、\_\_down函数
- 调用\_\_down\_common, 传入参量state = TASK\_UNINTERRUPTIBLE, 不可被信号申请中断:
- \_\_down\_common(sem, TASK\_UNINTERRUPTIBLE, MAX\_SCHEDULE\_TIMEOUT);
- MAX\_SCHEDULE\_TIMEOUT 表示无限期睡眠

```
static noinline void __sched __down(struct semaphore *sem)
{
    __down_common(sem, TASK_UNINTERRUPTIBLE, MAX_SCHEDULE_TIMEOUT);
}
```

### 问题三：简要分析down和up操作的实现

- **up操作的实现**

- 1、\_\_up函数
- 唤醒最早等待的沉睡进程，并从等待链表中删除此节点
- 2、up函数
- 上锁：raw\_spin\_lock\_irqsave(&sem->lock, flags);
- 等待队列为空，信号量加1：if (likely(list\_empty(&sem->wait\_list)))
- sem->count++;
- 等待队列非空，唤醒：else
- \_\_up(sem);
- 释放锁：raw\_spin\_unlock\_irqrestore(&sem->lock, flags);

```
static ninline void __sched __up(struct semaphore *sem)
{
    struct semaphore_waiter *waiter = list_first_entry(&sem->wait_list,
                                                         struct semaphore_waiter, list);

    list_del(&waiter->list);
    waiter->up = true;
    wake_up_process(waiter->task);
}
```

```
void up(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(list_empty(&sem->wait_list)))
        sem->count++;
    else
        __up(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}

EXPORT_SYMBOL(up);
```

## 问题三：简要分析down和up操作的实现

### • down操作的实现

#### • 1、\_\_down\_common函数

- (1) 创建一个等待进程的表并接到信号量结构体的等待进程表尾：
- list\_add\_tail(&waiter.list, &sem->wait\_list);
- (2) 将up设置为false，当进程被正确唤醒时up = true: waiter.up = false;
- (3) 休眠：for (;;)
- (4) 根据state参量的设置，被信号中断：if (signal\_pending\_state(state, current))
- goto interrupted;
- (5) 超时中断：if (unlikely(timeout <= 0))
- goto timed\_out;
- (6) 让出CPU：timeout = schedule\_timeout(timeout);
- (7) 正常唤醒：if (waiter.up)
- return 0;
- (8) 超时中断处理：timed\_out:
- list\_del(&waiter.list);
- return -ETIME;
- (9) 信号打断处理：interrupted:
- list\_del(&waiter.list);
- return -EINTR;

```
tatic inline int __sched __down_common(struct semaphore *sem, long state,
                                       long timeout)
{
    struct semaphore_waiter waiter;

    list_add_tail(&waiter.list, &sem->wait_list);
    waiter.task = current;
    waiter.up = false;

    for (;;) {
        if (signal_pending_state(state, current))
            goto interrupted;
        if (unlikely(timeout <= 0))
            goto timed_out;
        __set_current_state(state);
        raw_spin_unlock_irq(&sem->lock);
        timeout = schedule_timeout(timeout);
        raw_spin_lock_irq(&sem->lock);
        if (waiter.up)
            return 0;
    }

timed_out:
    list_del(&waiter.list);
    return -ETIME;

interrupted:
    list_del(&waiter.list);
    return -EINTR;
}
```