

实例分析

舒博文

李永康

张裕

Fileread & Filewrite 三个参数

```
fileread(struct file *f, char *addr, int n)
```

```
filewrite(struct file *f, char *addr, int n)
```

- 1.文件描述符（代表一个文件）
- 2.Addr表示用于保存从file文件中读取的数据的内存地址
- 3.读取/写入的数据大小

F-type的判断

```
struct file {  
    enum { FD_NONE, FD_PIPE, FD_INODE } type;  
    int ref; // reference count  
    char readable;  
    char writable;  
    struct pipe *pipe;  
    struct inode *ip;  
    uint off;  
};
```

F->type 的作用是用于表示文件类型

FD_NONE

FD_PIPE

FD_INODE

管道类型：管道是一个小的内核缓冲区，以文件描述符对的形式提供给进程

inode类型

文件的权限管理如何实现？

```
struct file {  
    enum { FD_NONE, FD_PIPE, FD_INODE } type;  
    int ref; // reference count  
    char readable;  
    char writable;  
    struct pipe *pipe;  
    struct inode *ip;  
    uint off;  
};
```

File结构体中存在名为readable和writable的域，分别表示可读和可写。

```
fileread(struct file *f, char *addr, int n)  
{  
    int r;  
  
    if(f->readable == 0)  
        return -1;
```

在fileread和filewrite函数中通过判断readable和writable是否为0来实现读和写的权限管理

```
filewrite(struct file *f, char *addr, int n)  
{  
    int r;  
  
    if(f->writable == 0)  
        return -1;
```

F->off变量代表什么？ 如何变化

- F->off变量代表相当于文件头的偏移量

```
if((r = readi(f->ip, addr, f->off, n)) > 0)
    f->off += r;
```

Readi成功后，使得off增加r，r表示的是从文件中读取数据的长度。

Readi&writei函数的参数ip类型是什么？ 包含的成员变量都是什么含义

```
readi(struct inode *ip, char *dst, uint off, uint n)
```

```
writei(struct inode *ip, char *src, uint off, uint n)
```

```
// in-memory copy of an inode
struct inode {
    uint dev;           // Device number
    uint inum;          // Inode number
    int ref;            // Reference count
    struct sleeplock lock; // protects everything below here
    int valid;          // inode has been read from disk?

    short type;         // copy of disk inode
    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+1];
};
```

Inode表示内核中的inode

相应的存在磁盘中的disk inode

Inode成员变量

```
// in-memory copy of an inode
struct inode {
    uint dev;           // Device number
    uint inum;          // Inode number
    int ref;            // Reference count
    struct sleeplock lock; // protects everything below here
    int valid;          // inode has been read from disk?

    short type;         // copy of disk inode
    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+1];
};
```

```
// On-disk inode structure
struct dinode {
    short type;         // File type
    short major;        // Major device number (T_DEV only)
    short minor;        // Minor device number (T_DEV only)
    short nlink;        // Number of links to inode in file system
    uint size;          // Size of file (bytes)
    uint addrs[NDIRECT+1]; // Data block addresses
};
```

- 1.Dev 设备号
- 2.Inum inode号
- 3.Ref 引用该inode的数量
- 4.Lock 锁，用于保护该inode
- 5.Valid 是否有效（从磁盘中读取）
//与disk inode中的成员变量相同
- 6.Type 文件类型
- 7.Major 最大设备数目
- 8.Minor 最小设备数目
- 9.Nlink 文件系统中该inode的链接数量
10. Size 文件大小
11. Addrs 数据块的地址

Readi函数中文件操作实际执行到的内容

```
int  
readi(struct inode *ip, char *dst, uint off, uint n)  
{  
    uint tot, m;  
    struct buf *bp;
```

```
    if(ip->type == T_DEV){  
        if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)  
            return -1;  
        return devsw[ip->major].read(ip, dst, n);  
    }
```

```
    if(off > ip->size || off + n < off)  
        return -1;
```

```
    if(off + n > ip->size)  
        n = ip->size - off;
```

```
    for(tot=0; tot<n; tot+=m, off+=m, dst+=m){  
        bp = bread(ip->dev, bmap(ip, off/BSIZE));  
        m = min(n - tot, BSIZE - off%BSIZE);  
        memmove(dst, bp->data + off%BSIZE, m);  
        brelse(bp);  
    }  
    return n;  
}
```

```
#define NDEV 10 // maximum major device number
```

```
#define T_DIR 1 // Directory  
#define T_FILE 2 // File  
#define T_DEV 3 // Device
```

处理类型为device类型的inode

如果出错，则返回-1

如果正确读出，返回devsw中读取的read值

当偏移量大于size或者n小于0时报错

当要读取的数据大小和偏移量的和大于size时重置n

从磁盘中读数据，存储到dst位置的内存中

```
#define BSIZE 512 // block size
```



```

writei(struct inode *ip, char *src, uint off, uint n)
{
    uint tot, m;
    struct buf *bp;

    if(ip->type == T_DEV){
        if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
            return -1;
        return devsw[ip->major].write(ip, src, n);
    }

    if(off > ip->size || off + n < off)
        return -1;
    if(off + n > MAXFILE*BSIZE)
        return -1;

    for(tot=0; tot<n; tot+=m, off+=m, src+=m){
        bp = bread(ip->dev, bmap(ip, off/BSIZE));
        m = min(n - tot, BSIZE - off%BSIZE);
        memmove(bp->data + off%BSIZE, src, m);
        log_write(bp);
        brelse(bp);
    }

    if(n > 0 && off > ip->size){
        ip->size = off;
        iupdate(ip);
    }

    return n;
}

```

Log_write写日志

更新内存中的inode（更新size）
更新磁盘中的inode（将内存的inode内容更新到磁盘对应的inode）

Bget函数的功能

```
bget(uint dev, uint blockno)
{
    struct buf *b;

    acquire(&bcache.lock);

    // Is the block already cached?
    for(b = bcache.head.next; b != &bcache.head; b = b->next){
        if(b->dev == dev && b->blockno == blockno){
            b->refcnt++;
            release(&bcache.lock);
            acquiresleep(&b->lock);
            return b;
        }
    }

    // Not cached; recycle an unused buffer.
    // Even if refcnt==0, B_DIRTY indicates a buffer is in use
    // because log.c has modified it but not yet committed it.
    for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
        if(b->refcnt == 0 && (b->flags & B_DIRTY) == 0) {
            b->dev = dev;
            b->blockno = blockno;
            b->flags = 0;
            b->refcnt = 1;
            release(&bcache.lock);
            acquiresleep(&b->lock);
            return b;
        }
    }
    panic("bget: no buffers");
}
```

```
struct {
    struct spinlock lock;
    struct buf buf[NBUF];

    // Linked list of all buffers, through prev/next.
    // head.next is most recently used.
    struct buf head;
} bcache;
```

```
struct buf {
    int flags;
    uint dev;
    uint blockno;
    struct sleeplock lock;
    uint refcnt;
    struct buf *prev; // LRU cache list
    struct buf *next;
    struct buf *qnext; // disk queue
    uchar data[BSIZE];
};
```

功能：输入设备号，块号。
先查找整个bcache，如果该block已经有cache与其对应，
则返回对应的cache
否则找到一块未被使用的cache返回。

结构体bcache用于管理所有的cache块，通过next和prev
指针形成了环状的结构。Head.next代表最近常用的
cache块。

算法：

使用的算法为LRU

通过head.prev和prev指针来寻找最近最不常用的cache
块。

Buffer cache中的dirty块是由哪个函数写入磁盘的?

```
bget(uint dev, uint blockno)
{
    struct buf *b;

    acquire(&bcache.lock);

    // Is the block already cached?
    for(b = bcache.head.next; b != &bcache.head; b = b->next){
        if(b->dev == dev && b->blockno == blockno){
            b->refcnt++;
            release(&bcache.lock);
            acquiresleep(&b->lock);
            return b;
        }
    }

    // Not cached; recycle an unused buffer.
    // Even if refcnt==0, B_DIRTY indicates a buffer is in use
    // because log.c has modified it but not yet committed it.
    for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
        if(b->refcnt == 0 && (b->flags & B_DIRTY) == 0) {
            b->dev = dev;
            b->blockno = blockno;
            b->flags = 0;
            b->refcnt = 1;
            release(&bcache.lock);
            acquiresleep(&b->lock);
            return b;
        }
    }
    panic("bget: no buffers");
}
```

```
struct buf*
bread(uint dev, uint blockno)
{
    struct buf *b;

    b = bget(dev, blockno);
    if((b->flags & B_VALID) == 0) {
        iderw(b);
    }
    return b;
}
```

```
// Sync buf with disk.
// If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
// Else if B_VALID is not set, read buf from disk, set B_VALID.
void
iderw(struct buf *b)
{
    struct buf **pp;

    if(!holdingsleep(&b->lock))
        panic("iderw: buf not locked");
    if((b->flags & (B_VALID/B_DIRTY)) == B_VALID)
        panic("iderw: nothing to do");
    if(b->dev != 0 && !havedisk1)
        panic("iderw: ide disk 1 not present");

    acquire(&idelock); //DOC:acquire-lock

    // Append b to idequeue.
    b->qnext = 0;
    for(pp=&idequeue; *pp; pp=&(*pp)->qnext) //DOC:insert-queue
        ;
    *pp = b;

    // Start disk if necessary.
    if(idequeue == b)
        idestart(b);

    // Wait for request to finish.
    while((b->flags & (B_VALID/B_DIRTY)) != B_VALID){
        sleep(b, &idelock);
    }

    release(&idelock);
}
```

linux 内核代码中 buffer cache 的替换是如何实现的？ 替换算法和 xv6 有何不同

Linux内核中文件Cache替换的具体过程是这样的：刚刚分配的Cache项链入到active_list头部，并将其状态设置为active，当内存不够需要回收Cache时，系统首先从尾部开始反向扫描active_list并将状态不是referenced的项链入到inactive_list的头部，然后系统反向扫描inactive_list，如果所扫描的项的处于合适的状态就回收该项，直到回收了足够数目的Cache项。

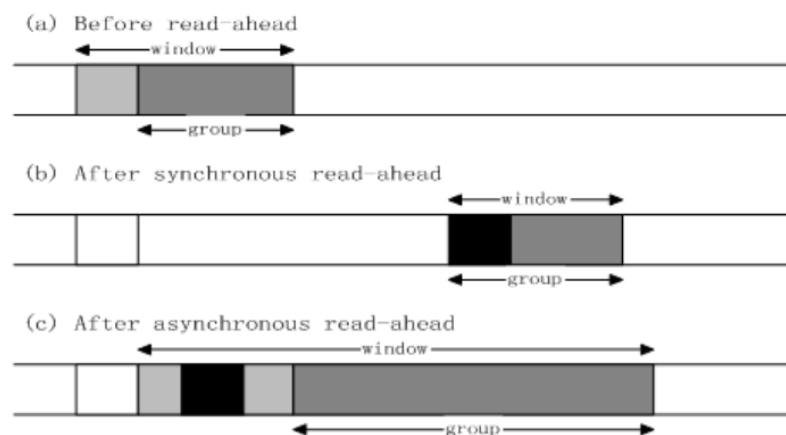


图 5 Linux 文件预读机制示意图

Linux中存在两个队列。替换算法也为LRU
与xv6不同的地方在于
Xv6为环状的形式，通过反向查找的方式来实现替换最长时间没有使用的页